

Deep Generative Models III

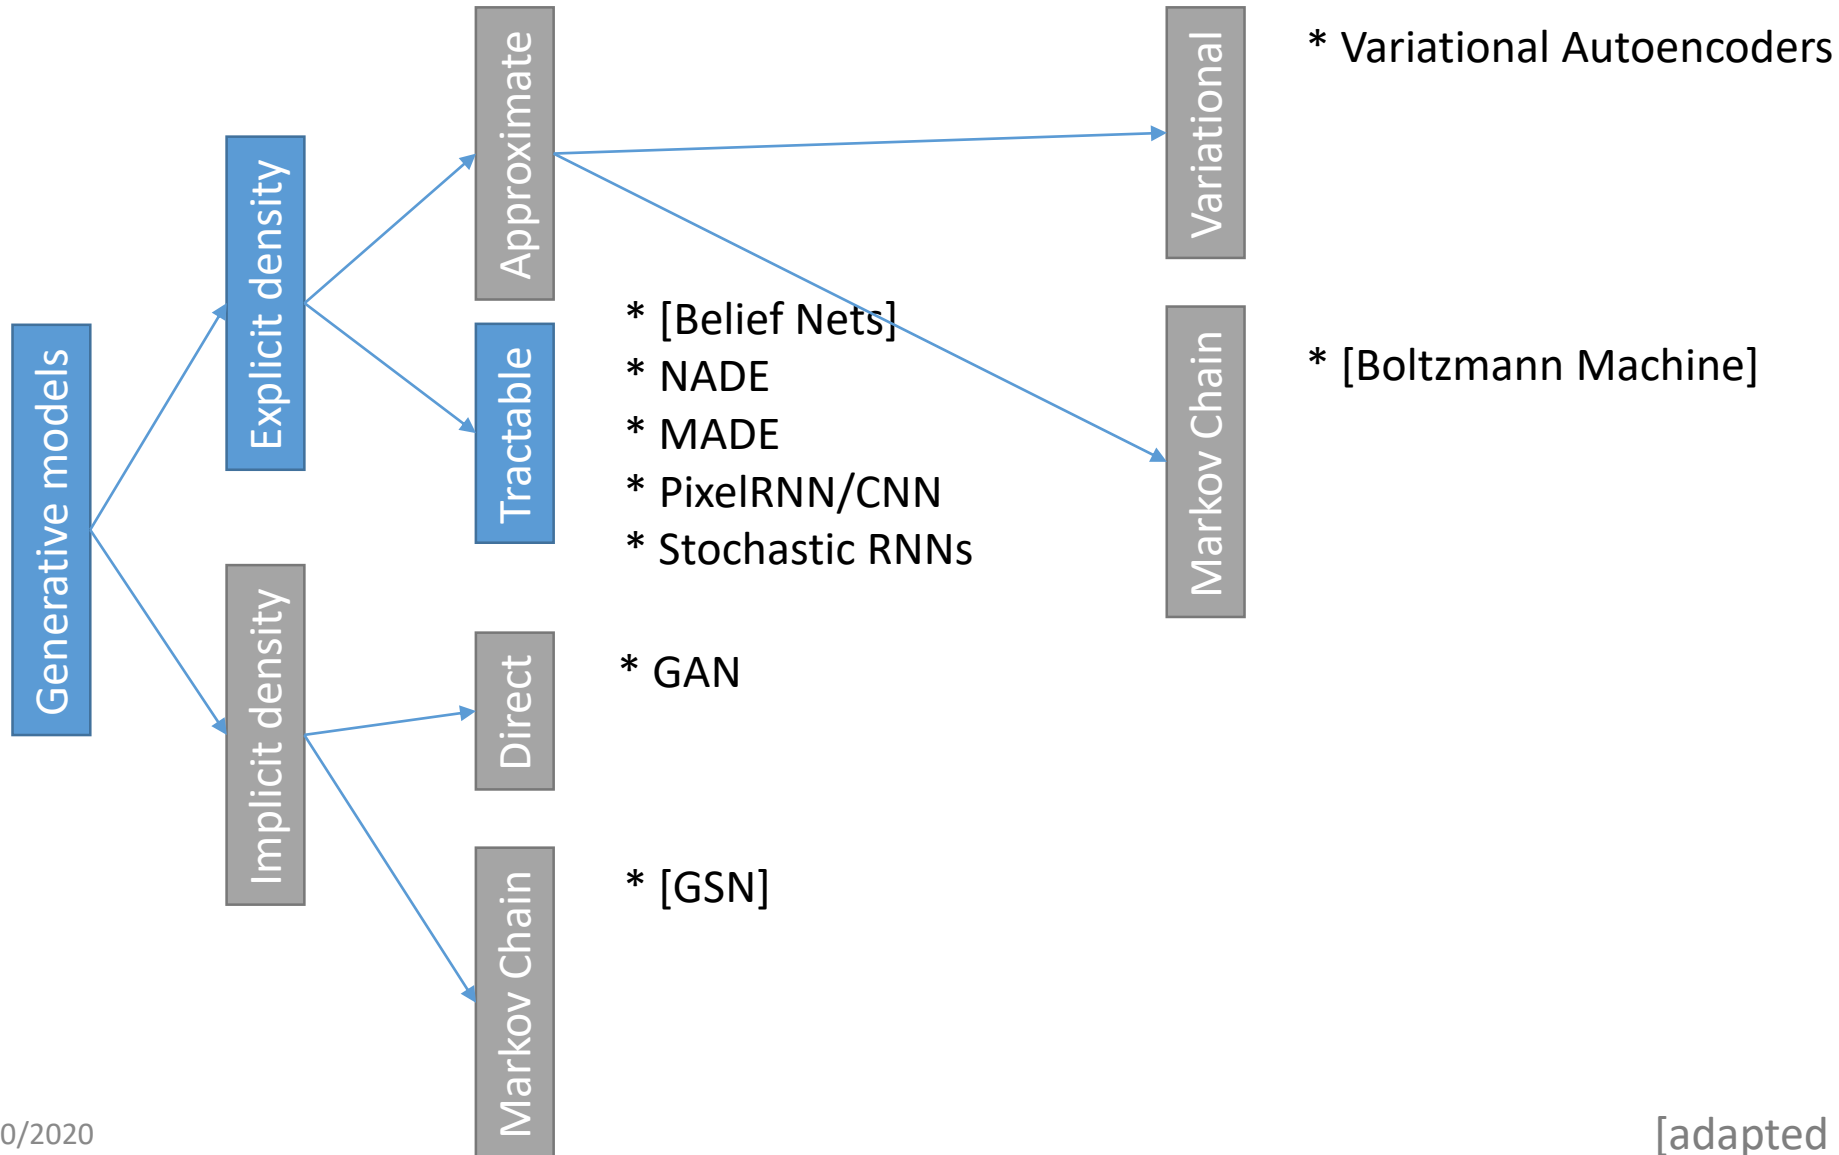
Autoregressive Models – Part I

Machine Perception

Otmar Hilliges

30 April 2020

Taxonomy of Generative Models



Auto-regressive property

A regression model, such as linear regression, models an output value based on a linear combination of input values. For example:

$$\hat{y} = b_0 + b_1 x_1$$

This technique can be used on time series where input variables are taken as observations at previous time steps, called lag variables.

For example, we can predict the value for the next time step (t+1) given the observations at the last two time steps (t-1 and t-2). As a regression model, this would look as follows:

$$x_t = b_0 + b_1 x_{t-1} + b_2 x_{t-2}$$

Because the regression model uses data from the same input variable at previous time steps, it is referred to as an autoregression (regression of self).

Sequence models are generative models

Discriminative models model the conditional distribution $P(Y|X)$

A generative model models the joint distribution $P(X, Y)$ of the observation X and the target Y or $P(X)$ in the unsupervised setting (density estimation)

A sequence model deals with sequential data:

- mapping sequences to scalars (e.g. language models)
- mapping seq. to seq. (e.g. machine translation models)

Sequence models are generative:

- Given seed x_1, x_2, \dots, x_k predict x_{k+1} .
- Then use x_2, x_3, \dots, x_{k+1} to predict x_{k+2}
- rinse & repeat

Learning the distribution of natural data

$$p(\mathbf{x}) = \prod_k \prod_j \prod_i p(\mathbf{x}_{i,j,k} | \mathbf{x}_{<})$$

$\{x_1, \dots, x_{i-1}\}$

PixelRNN/PixelCNN (Images)

[van den Oord, Kalchbrenner, Kavukcuoglu, 2016]

[van den Oord, Kalchbrenner, Vinyals, et al. 2016]

Video Pixel Nets (Videos)

[Kalchbrenner, van den Oord, Simonyan, et al. 2016]

ByteNet (Language/Seq2Seq)

[Kalchbrenner, Espeholt, Simonyan, et al. 2016]

WaveNet (Audio)

[van den Oord, Dieleman, Zen, et al. 2016]

Attempt 1: Autoregressive models – tabular approach

Via the chain rule of probabilities we can factorize the joint distribution over the n-dimensions:

$$p(x) = \prod_1^n p(x_i | x_1, \dots, x_{i-1}) = \prod_1^n p(x_i | \underline{x_{\leq i}})$$

param;

$$2^{n-1} - 1$$

↑
binary variables



$1, \dots, i-1$

given some
ordering of \mathcal{D}


Attempt 2: Autoregressive generative models

Specify conditionals with fixed number of parameters:

Assume $p(x_i | x_{<i})$ to correspond to Bernoulli random variable and learn to map previous inputs to the mean: 期望值 p

$$p_{\theta_i}(x_i | x_{<i}) = \text{Ber}(f(\underbrace{x_1, \dots, x_{i-1}}))$$

$$\sum_{i=1}^n |\theta_i| = O(n^2)$$




Fully Visible Belief Networks

$$\hat{x}_i = p(x_i = 1 | x_{<i})$$

Modelled via logistic regression:

$$f_i(x_1, x_2, \dots, x_{i-1}) = \sigma(\underbrace{\alpha_0^i}_{\text{bias}} + \underbrace{\alpha_1^i}_{\text{weight}} x_1 + \dots + \underbrace{\alpha_i^i}_{\text{weight}} x_{i-1})$$

$$\# \sum_i^n |\alpha_i| = \mathcal{O}(n^2)$$



Neural Autoregressive Density Estimator (NADE)

autoencoder-like neural network to learn

$p(x_i = 1 | x_{<i})$: \leftarrow binary data

$$\hat{h}_i = \sigma(b + W_{\cdot, <i} x_{<i})$$

first $i-1$ col


$$\hat{x}_i = \sigma(c_i + V_{i, \cdot} h^{(i)})$$

ith row


each conditional is modeled by the same neural network

We can leverage the fact that:

$$(b + W_{\cdot, <i+1} x_{<i+1}) - (b + W_{\cdot, <i} x_{<i}) = W_{\cdot, i+1} x_{i+1}$$



NADE schematically



[image source: A. Courville. UdM]

variants of NADE could work on non-binary data

NADE – Training

training by maximizing the average log-likelihood:

$$\frac{1}{T} \sum_{t=1}^T \log(p(x^t)) = \frac{1}{T} \sum_{t=1}^T \sum_{i=1}^D \log p(x_i^{(t)} | x_{<i}^{(t)})$$

Advantages

- efficient: computations are in $O(TD)$
- could make use of second-order optimizers
- easily extendable to other types of observations (reals, multinomials)

Paper explores different orderings of inputs

- random order works fine!

Next Week

Wrap-up deep generative modelling

Reinforcement Learning Bootcamp