

Reinforcement Learning

Bootcamp

Machine Perception

7 May 2020

Overview

- Motivation
- Dynamic Programming and Planning
- Temporal Difference Learning

Overview Learning Methods

Supervised Learning

- Learn mapping between input and output
- Labelled ground-truth data guides the learning process



Unsupervised Learning

- Learn a structure in a dataset
- No ground-truth data available



[SR-GAN. Ledig et al. CVPR '17]

Reinforcement Learning

- Learn how to act
- Data revealed through interaction with an environment





Image source: pathmind.com


Reinforcement Learning Problem Statement

Reinforcement Learning is a problem, not a method:


- Learning how to map states to actions
- Maximize a numerical reward signal
- Unknown and uncertain environment



Motivation




Games



Robot Control

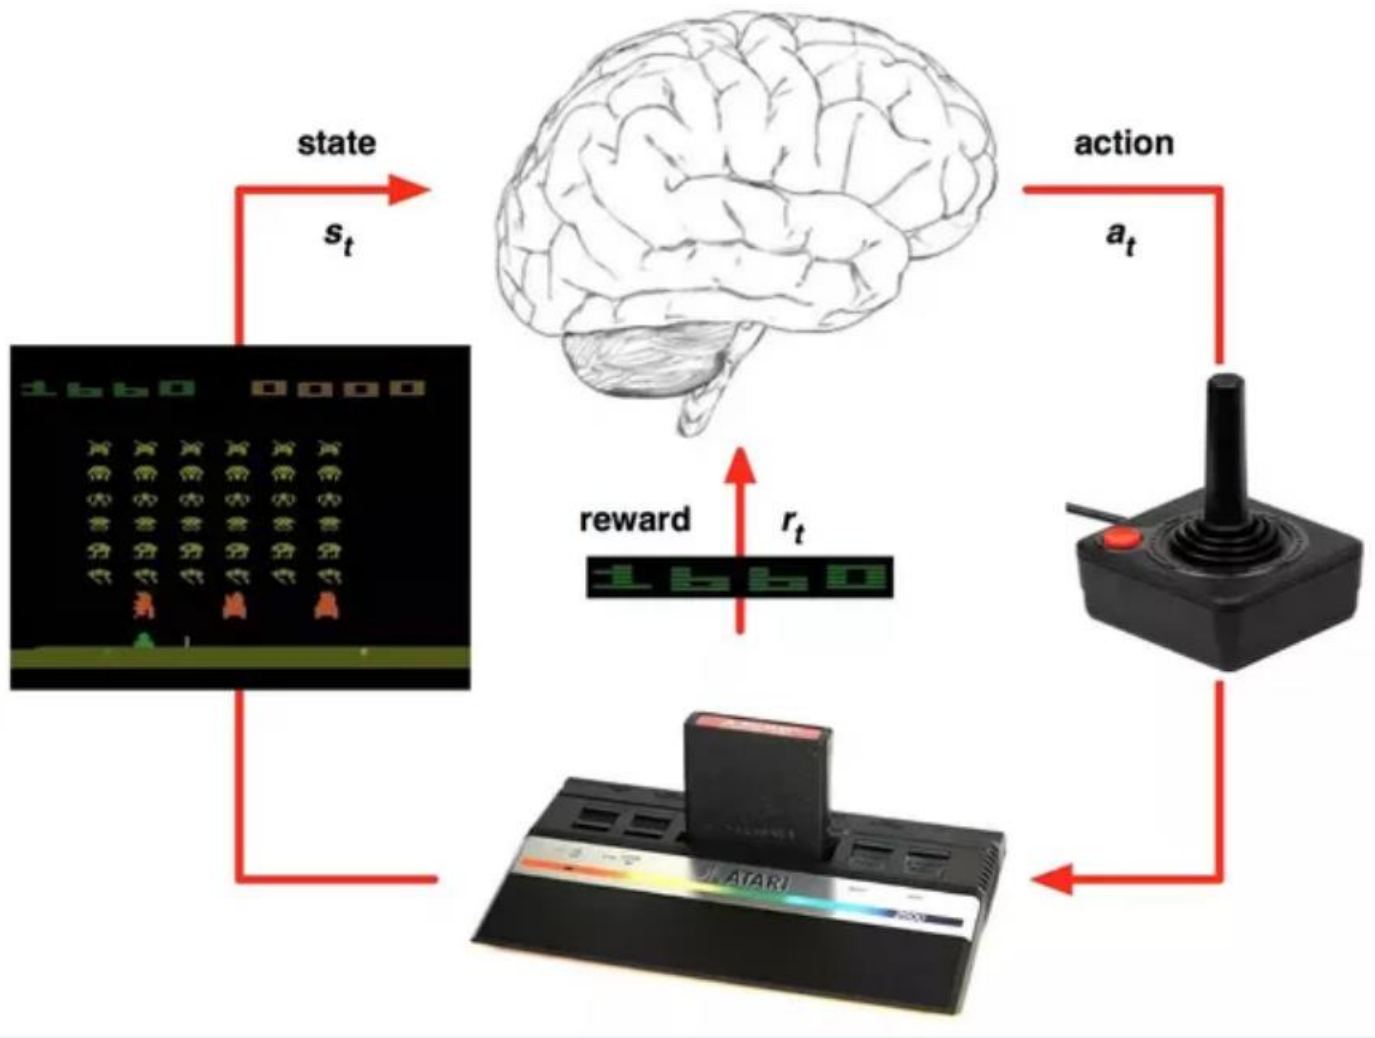


Attention based models



Logistics and Operations

Example: Atari



Objective: Achieve a high score


State: Images of the game state

Actions: Move the joystick


Reward: Score of the game

[Image source: RLSS summer school – Olivier Pietquin]

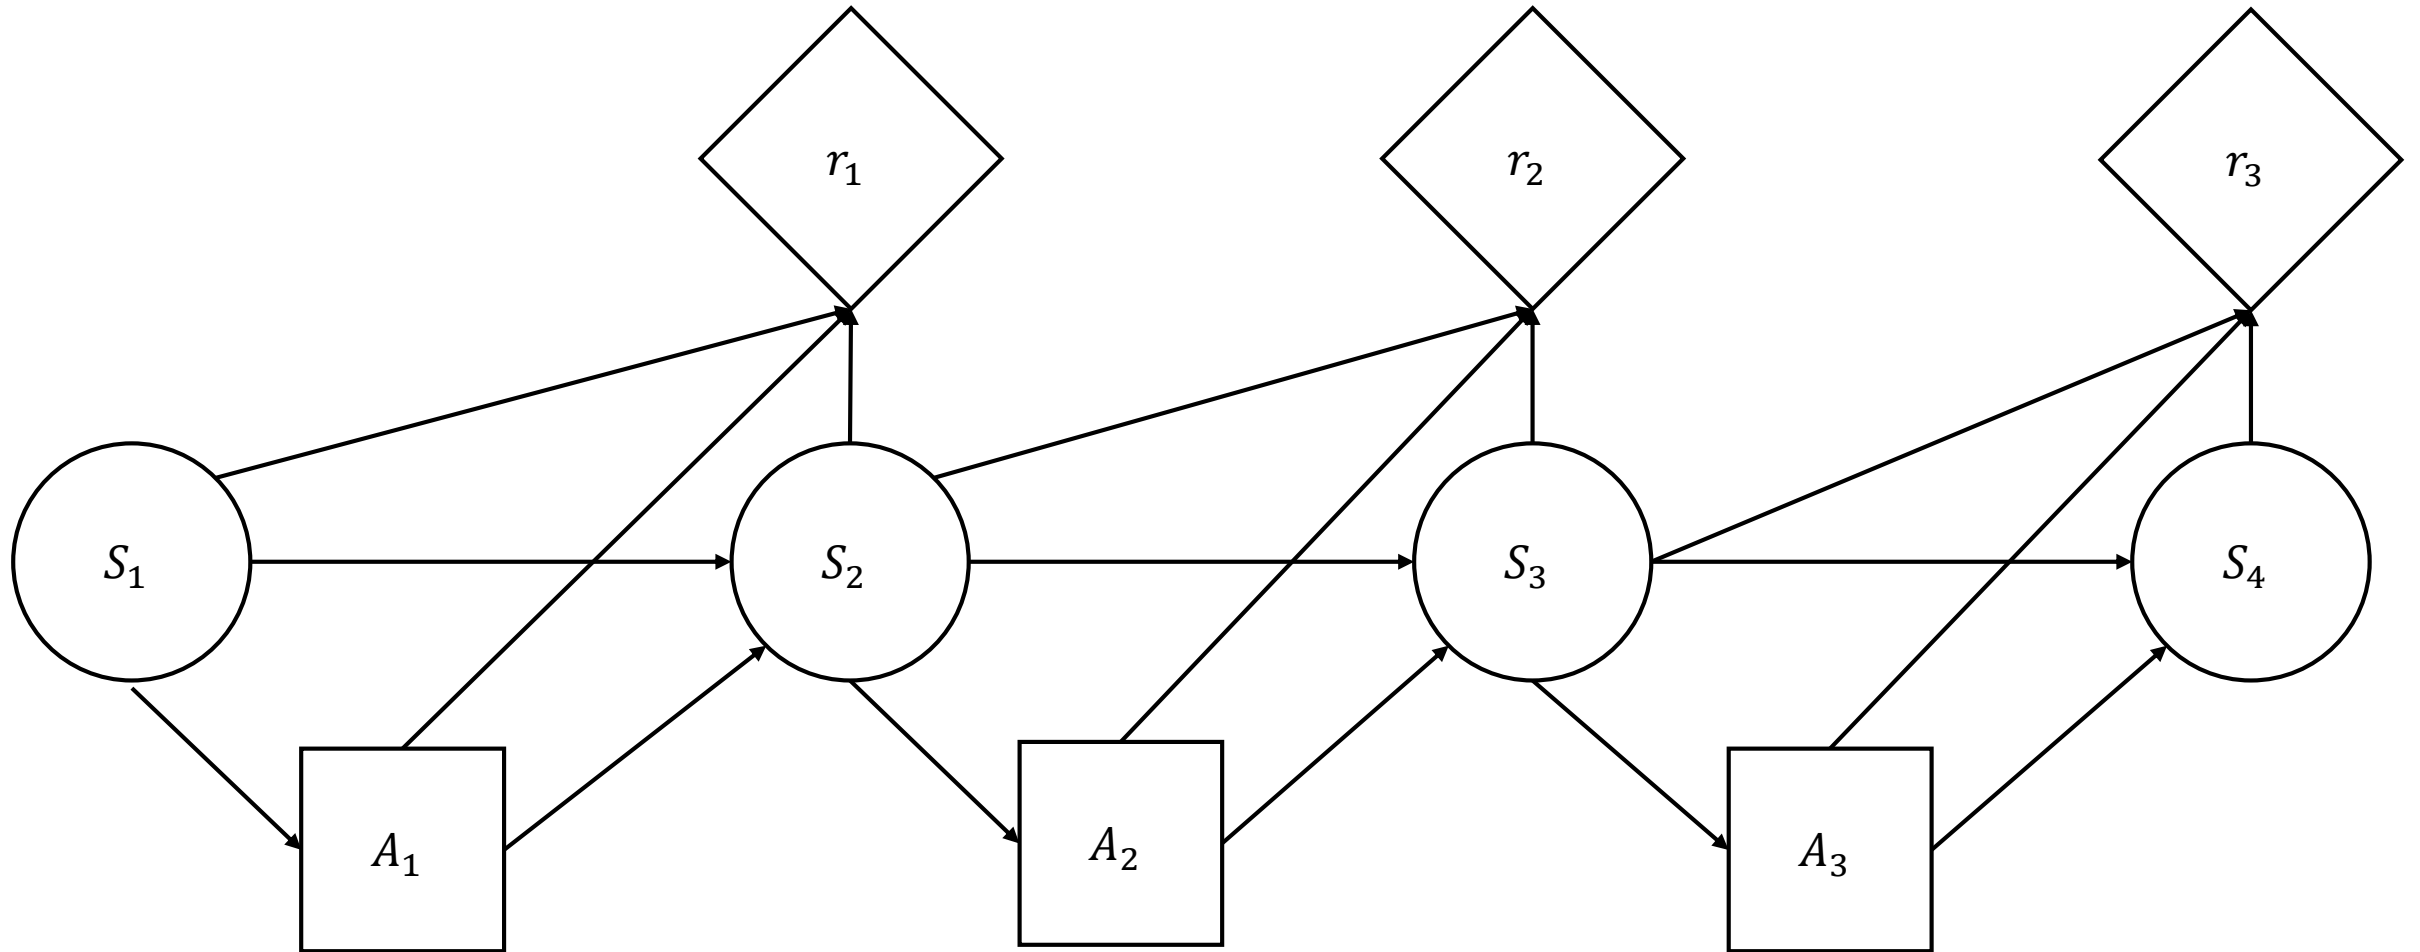
Recap Markov Decision Processes



Recap Markov Decision Processes



Recap Markov Decision Processes



Recap Markov Decision Processes

We model the environment as a Markov Decision Process

- Assumption: Transitions, initial state, rewards and policies are deterministic \rightarrow Simpler notation.

An MDP consists of $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$

- a set of states \mathcal{S}
- a set of actions \mathcal{A}
- a reward function $r: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ (deterministic, can also be a distribution)
- a transition function $p: \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ (deterministic, can also be a distribution)
- initial state $s_0 \in \mathcal{S}$
- Discount factor: γ

Problem

The core problem in solving MDPs is to find a “policy” $\pi(s)$ for the decision maker

- Once we know the optimal policy, for each state of the decision process we know how to act
- The chosen action is fully determined by $\pi(s)$

The optimal policy will maximize some cumulative quantity over the random rewards

Value Function

Important concept: Value function $V_\pi: \mathcal{S} \rightarrow \mathbb{R}$ of a policy π .

V_π expresses the expected cumulative reward (return) for each state that we achieve when following the policy π

Bellman equation for v_π

$$G_{t+1} = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} = \gamma^0 R_{t+1} + \sum_{k=1}^{\infty} \gamma^k R_{t+k+1}$$

$$= R_{t+1} + \gamma \sum_{k=1}^{\infty} \gamma^{k-1} R_{t+k+1}$$

$$= R_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k R_{t+(k+1)+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= R_{t+1} + \gamma G_{t+1}$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

Probability of taking an action given a state

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r \boxed{p(s', r | s, a)} [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

Joint probability of next state and the reward given the current state and action.

Transition Matrix

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

The expected reward given a new state

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \left[\sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']] \right]$$

A weighted sum of rewards given the current state and action

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

A weighted sum of rewards given the current state

Bellman equation for v_π

$$G_t \doteq \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$v_\pi(s) \doteq E_\pi[G_t | S_t = s]$$

$$= E_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) [r + \gamma E_\pi[G_{t+1} | S_{t+1} = s']]$$

$$= \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$

Next Steps

How can we solve the Bellman Equation?

- **Dynamic Programming**

- Requires MDP with given environment model (transition probabilities)
- Value function based on estimates of subsequent values

$p(s', r | s, a)$



- **Monte-Carlo Methods** (not covered here)

- No need to known environment model
- Value function based on gathered data and only updated along full trajectories

- **Temporal-Difference Learning**

- No need to known environment model
- Combination of DP and MC principles
- Value function updated based on estimates of subsequent values along trajectory

Dynamic Programming

Overview

- Collection of algorithms that can compute optimal policies given a perfect model of the world (MDP).
- Can be applied to continuous space, we focus on discrete.
- Limited utility, great theoretical importance.
 - All RL methods have a similar goal, with less computation and no perfect model assumption.

Overview – 3 Key ideas

Value Iteration for estimating π_*

- 1) Compute optimal state value function v_*
- 2) Obtain policy π_* based on v_*

*4/ initialize V_0
iterate V until converge
then we get π^**

Iterative Policy Evaluation

- 1) Compute the state value function v_π for a given policy π

Policy Iteration for estimating π_*

- 1) Compute the state value function v_π for any arbitrary policy π
- 2) Update policy π given $v_\pi \rightarrow \pi'$
- 3) Iterate till $\pi \approx \pi'$

Overview – 3 Key ideas

Value Iteration for estimating π_*

- 1) Compute optimal state value function v_*
- 2) Obtain policy π_* based on v_*

Iterative Policy Evaluation

- 1) Compute the state value function v_π for a given policy π

Policy Iteration for estimating π_*

- 1) Compute the state value function v_π for any arbitrary policy π
- 2) Update policy π given $v_\pi \rightarrow \pi'$
- 3) Iterate till $\pi \approx \pi'$

Example Environment

0	0	0	Goal 100
0	0	Trap -100	Trap -100
Start	0	0	0

Goal: Find optimal path to goal

States: Grid fields
Start state blue field
Terminal states trap/goal

Actions: Move left/right/up/down

Reward: 100 upon reaching the goal
-100 landing in the trap
-5 for each step taken

Other info: $\gamma = 1$
Initial states initialized to 0
Deterministic transitions

Value Function

- Given a value function V_π we can derive a greedy policy π' :

$$\pi'(s) = \operatorname{argmax}_{\substack{a \in A \\ \text{red oval}}} (r(s, a) + \gamma V_\pi(p(s, a)))$$

- One can show that $V_{\pi'}(s) \geq V_\pi(s)$

Value Function

- For the optimal policy π^* : The greedy policy of V_{π^*} is π^* itself.
- V_{π^*} fulfills the Bellman Optimality Equation:

$$V^*(s) = \max_{a \in A} \{r(s, a) + \gamma V^*(p(s, a))\}$$

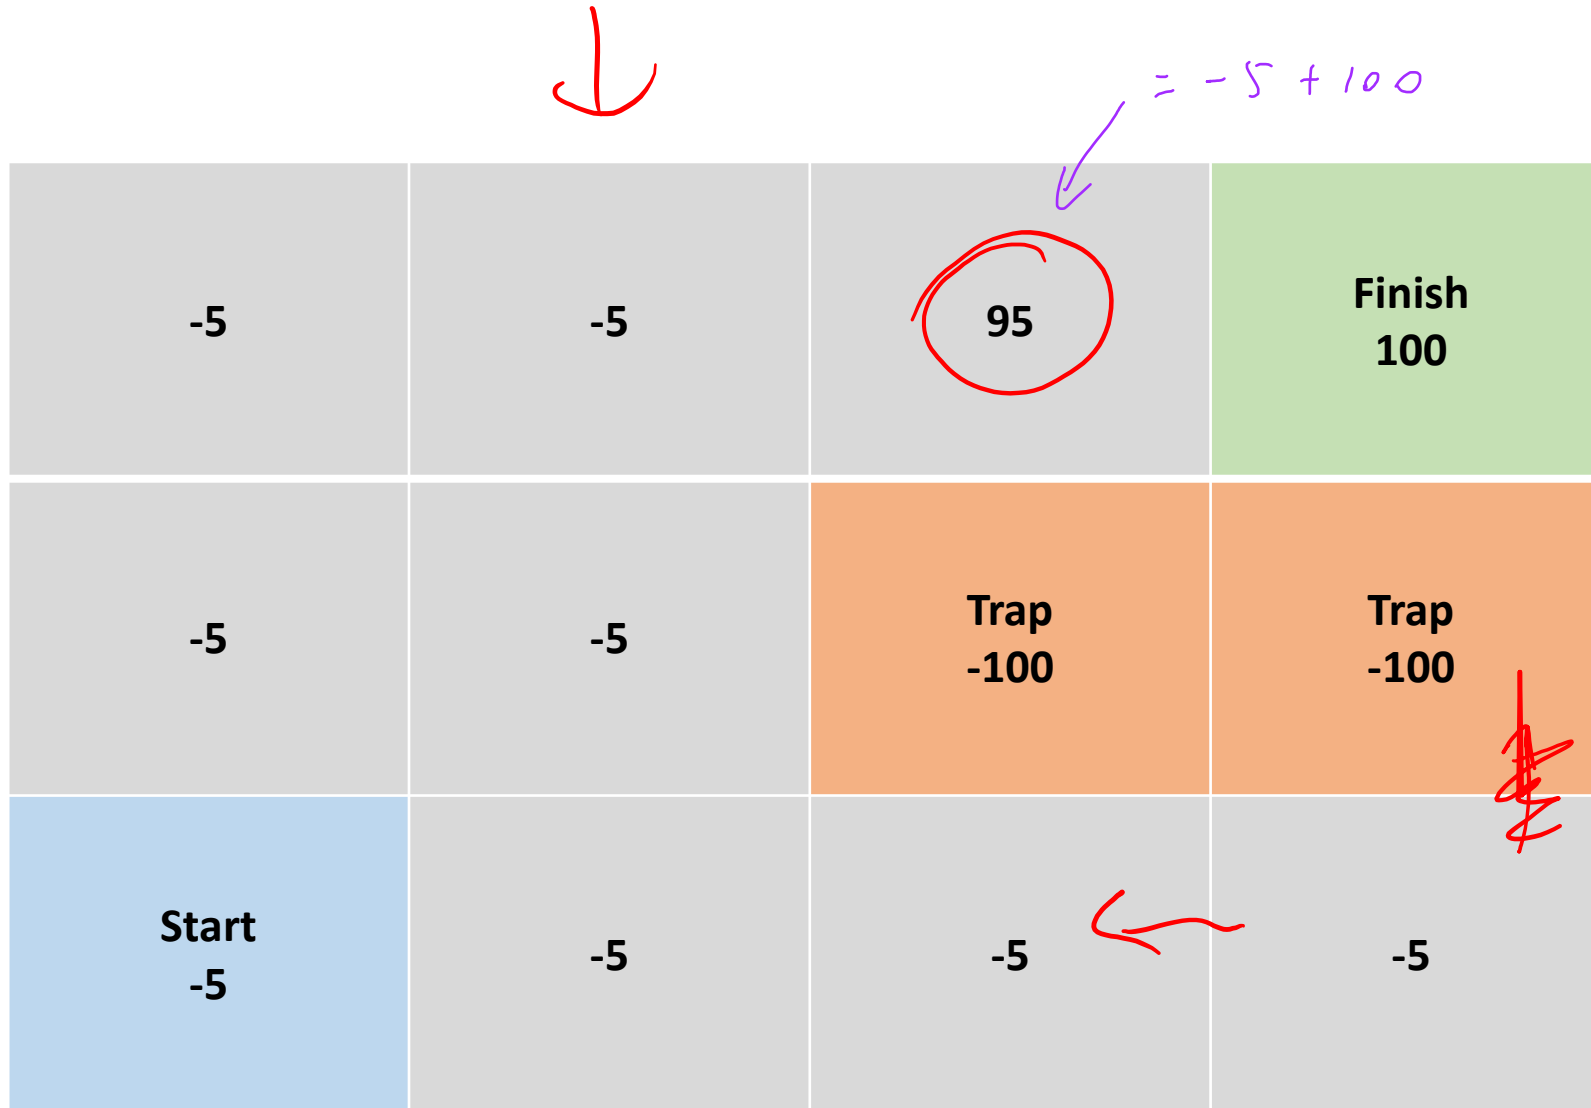
Example of Value Iteration ($\gamma = 1$, step cost=5)

Initial value is 0 for every state

0	0	0	Goal 100
0	0	Trap -100	Trap -100
Start 0	0	0	0

Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 1



Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 2

-10	90	95	Goal 100
-10	-10	Trap -100	Trap -100
Start -10	-10	-10	-10

Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 3

85	90	95	Goal 100
-15	85	Trap -100	Trap -100
Start -15	-15	-15	-15

Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 4

85	90	95	Goal 100
80	85	Trap -100	Trap -100
Start -20	80	-20	-20

Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 4

85	90	95	Goal 100
80	85	Trap -100	Trap -100
Start 75	80	75	-25

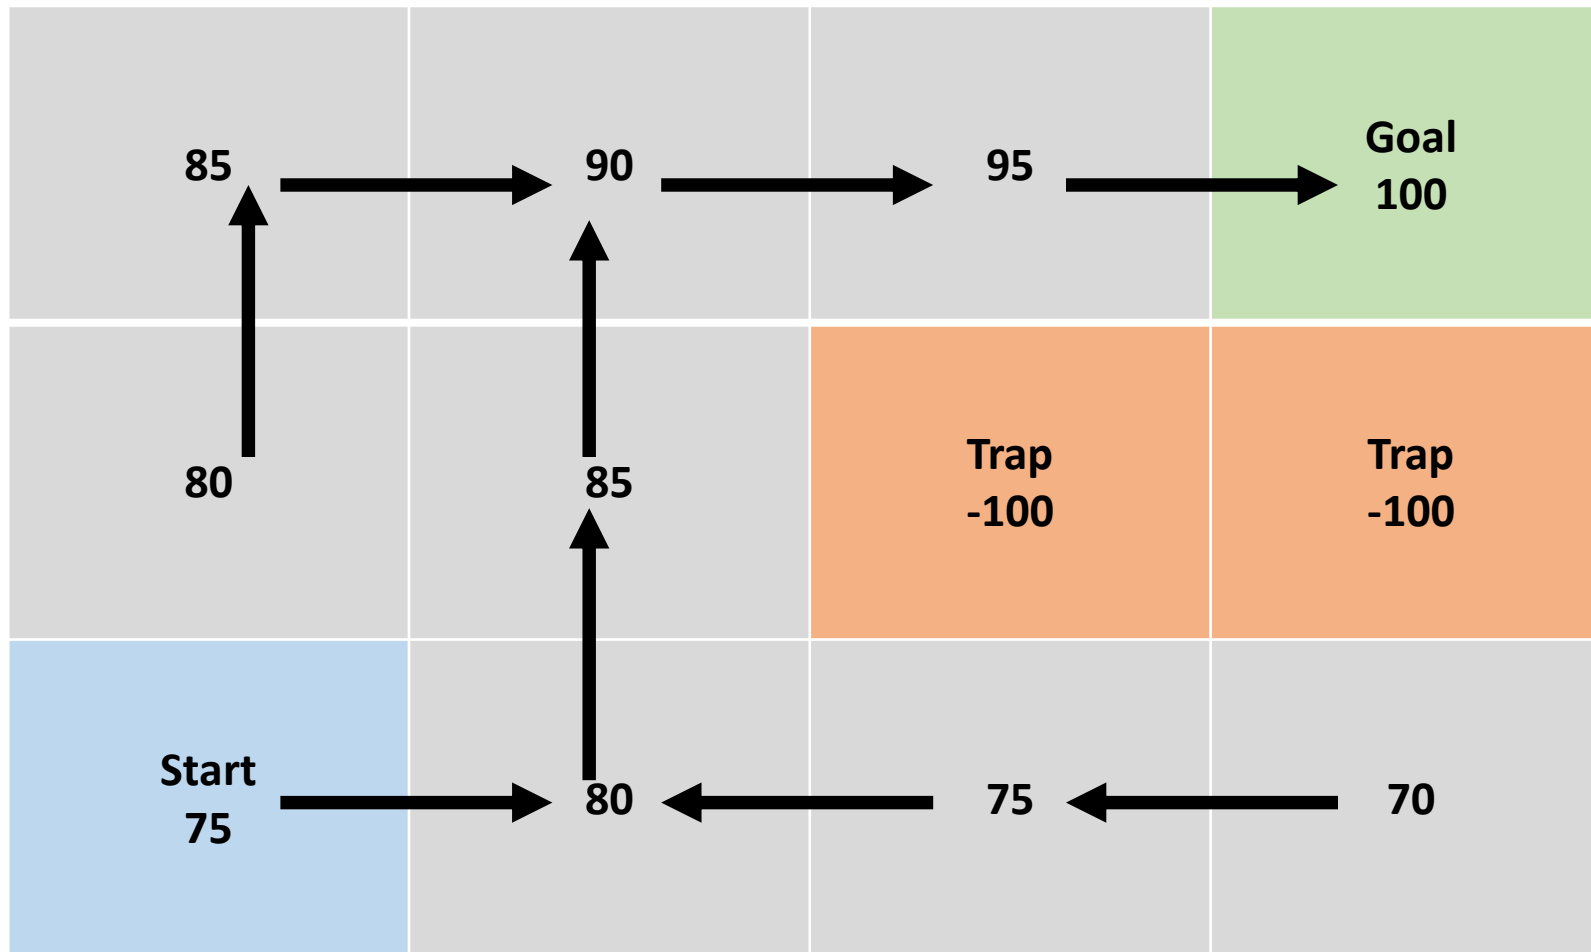
Example of Value Iteration ($\gamma = 1$, step cost=5)

Iteration 5. Done as there is no change anymore.

85	90	95	Goal 100
80	85	Trap -100	Trap -100
Start 75	80	75	70

Example of Value Iteration ($\gamma = 1$, step cost=5)

Compute greedy policy




Value Iteration: Step-wise

1. $V_0^*(s)$ = optimal value for state s when $H=0$

- $\underline{V_0^*(s)} = 0 \forall s$

2. $V_1^*(s)$ = optimal value for state s when $H=1$

- $V_1^*(s) = \max_{a \in A} \sum_{s', r} p(s', r | s, a) [r + \gamma \underline{V_0^*(s')}]$



3. $V_2^*(s)$ = optimal value for state s when $H=2$

- $V_2^*(s) = \max_{a \in A} \sum_{s', r} p(s', r | s, a) [r + \gamma V_1^*(s')]$

4. $V_k^*(s)$ = optimal value for state s when $H=k$

- $V_k^*(s) = \max_{a \in A} \sum_{s', r} p(s', r | s, a) [r + \gamma V_{k-1}^*(s')]$

Value Iteration: Algorithm


Instead of continuously doing policy evaluation and policy improvement, we can do the improvement step in each evaluation iteration.

```

$$\begin{aligned} & i \leftarrow 0 \\ & V_0 \leftarrow 0 \\ & \text{while } \Delta > \theta \{ \\ & \quad \Delta \leftarrow 0 \\ & \quad \text{for all } s \in S \{ \\ & \quad \quad v \leftarrow V(s) \\ & \quad \quad V(s) \leftarrow \max_{a \in A} \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')] \\ & \quad \quad \Delta \leftarrow \max(\Delta, |v - V(s)|) \\ & \quad \} \\ & \} \\ & \pi_i \leftarrow \text{greedy policy from } V_i \end{aligned}$$

```

always 1 or 0 in our case due to deterministic environment



Value Iteration vs. Policy Iteration Algorithm

Value Iteration:

1. Iterate until value function converges.
2. Obtain optimal policy through the greedy policy.

Policy Iteration:

1. Initialize arbitrary policy.

Repeat until policy does not improve:

2. Evaluate the value function given the current policy.
3. Improve the policy.

Pros and Cons of Dynamic Programming

Pros:

- Exact methods
- Policy/Value iteration are guaranteed to converge in finite number of iterations
- Value iteration is typically more efficient

Cons:

- Need to know the transition probability matrix
- Need to iterate over the whole state space (very expensive)
- Requires memory proportional to the size of the state space

Summary

This Lecture

Dynamic Programming

- Uses value of neighbouring states to update current state (given π)
- Visits every state
- Requires knowledge of the transition probabilities

Monte Carlo Methods (not covered in detail)

- Run in episodes, update state values at end of episode.
- Experience based (no need to know system dynamics)
- High Variance

Temporal Differences

- Run in episodes, update state values at each step.
- Combination of DP and MC ideas
- Exploration/Exploitation → Local minima

Next Lecture

- TD-Learning
- Function Approximation
 - Extend the ideas to arbitrary large state spaces (e.g. images).
- Policy based gradient methods
 - Replace action-value methods with parametrised policies.
- Deep Reinforcement Learning applications
- Research in Deep Reinforcement Learning

References

Book by Sutton:

<http://incompleteideas.net/book/the-book-2nd.html>

Lecture by David Silver:

<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>

Lecture by John Schulman:

https://www.youtube.com/watch?v=aUrX-rP_ss4